

One Pipeline, Any Device

Swap the Reader, Not the Pipeline

Tom Kaltofen, mloda.ai

One process, two implementations

Data engineering

batch pipeline
semantic layer

ML

offline feature
online feature


AI agents

index, embeddings
tool call

Robotics


log replay, sim
on-board

A full pipeline was the fastest way to create value



Two real runs, drawn from their traces. Nothing shared.

Copying a pipeline is free now. That is the problem.
One definition. Small readers around it.



The same two runs, one definition, two readers.

One process, one definition.

Data engineering

ML

AI agents

Robotics

nearest_ahead_m

Small readers around it, and a receipt
for every number.

mloda.



github.com/mloda-ai/mloda-demo

How far is 1200?

Kinect, Azure Kinect, RealSense D435, ROS 16UC1, ScanNet, Redwood	1.20 m
RealSense L515 (same vendor, 0.25 mm units)	0.30 m
TUM RGB-D (5000 per metre)	0.24 m
SUN RGB-D (bits rotated first)	0.15 m
KITTI depth (256 per metre)	4.69 m
ROS 32FC1 (float, metres)	1.2 km
NYU Depth raw (Kinect disparity)	not a distance
Depth Anything, MiDaS (relative depth)	no unit at all